




POLITECNICO
MILANO 1863

SCUOLA DI INGEGNERIA INDUSTRIALE
E DELL'INFORMAZIONE



Title

TESI DI LAUREA MAGISTRALE IN

XXXXXXXXXXXXX ENGINEERING - INGEGNERIA XXXXXXXXXXXXXXXX

Jose Luis Vilaseca Gallego, 11140177
Luis Alfonso Pérez Mijangos, 11133173

Advisor:

Prof. Name Surname

Co-advisors:

Name Surname

Name Surname

Abstract:

Here goes the text of the abstract

Academic year:

20XX-20XX

Key-words: here, the keywords, of your thesis

1. Introduction

Here goes the introduction

2. Problem Statement

3. Problem Formulation and Motivation

3.1. Problem Statement

The primary objective of this project is to estimate the 3D trajectory and spin dynamics of a bowling ball from a static monocular video feed for advanced sports analysis. Because intrinsic camera parameters are not provided, the system must rely on projective geometry and physical constraints.

To achieve a robust and accurate estimation, the problem is formulated as a structured computer vision pipeline consisting of five interdependent tasks:

1. **Lane Detection:** Identifying the physical boundaries of the playing surface within the image plane, despite perspective distortion and object occlusion.
2. **Planar Metric Rectification:** Utilizing the detected lane boundaries to compute a planar homography matrix, mapping the perspective-distorted image space to a metric, 2D bird's-eye view.
3. **Ball Detection and Tracking:** Isolating the bowling ball across consecutive frames and applying temporal interpolation to generate a continuous, robust 2D trajectory from discrete, raw center coordinates.

4. **3D Trajectory Estimation:** Lifting the 2D tracking data into 3D spatial coordinates by enforcing two geometric constraints: the ball maintains constant contact with the rectified lane plane, and it possesses a known, fixed physical radius.
5. **Spin Calculation:** Analyzing the rotational displacement of surface features on the ball to estimate its spin rate, a critical parameter in bowling performance.

3.2. Motivation and Relevance

In modern sports analytics, precise metrics such as ball speed, entry angle, trajectory curvature, and spin rate are fundamental for player evaluation. In bowling, the interaction between the ball’s spin and the lane’s oil pattern dictates the success of a shot.

Commercial systems capable of extracting these metrics typically rely on complex, synchronized multi-camera stereoscopic setups or specialized sensor arrays that are highly cost-prohibitive. Developing a solution that extracts both robust 3D trajectory and spin parameters from a single, uncalibrated static camera democratizes access to professional-grade analytics. It demonstrates the efficacy of combining modern object tracking with classic geometric priors to extract deep metric insights from standard video footage.

3.3. State of the Art

3.3.1 Lane Detection and Metric Rectification

The foundational techniques for boundary extraction are heavily derived from Advanced Driver Assistance Systems (ADAS). Standard pipelines in dynamic environments rely on localized gradient filters (such as the Sobel or Canny operators) and the Hough Transform to extract linear boundaries [?].

However, translating pixel-space detections into metric measurements traditionally requires full camera calibration [?]. In structured sports environments, where the playing surface is guaranteed to be a flat plane and the camera is static, the state of the art often bypasses explicit calibration in favor of Planar Homography [?]. By identifying four coplanar control points with known real-world dimensions (such as the corners of a tennis court or bowling lane), a homography matrix can be computed to rectify the image, providing a highly accurate metric mapping of the surface [?].

3.3.2 Ball Tracking, Trajectory Estimation, and Spin

Tracking fast-moving sports balls is challenging due to motion blur, occlusion, and background clutter. Modern solutions frequently utilize deep learning frameworks (e.g., TrackNet) or background subtraction to isolate the ball [?]. To transition from 2D pixel tracking to 3D spatial estimation without depth sensors, domain-specific applications enforce physical priors. By leveraging the geometric constraint that a rigid sphere maintains contact with a known ground plane, the object’s spatial coordinates can be robustly estimated [?]. Furthermore, estimating spin from a monocular feed is notoriously difficult. State-of-the-art approaches typically involve analyzing the frame-by-frame displacement of surface textures (such as finger holes, logos, or dotted patterns) using feature-matching algorithms, template matching, or optical flow [?] (AGREGAR CITAS OPTICAL FLOW) . Recent studies have even explored event-based cameras to capture ultra-fast spin dynamics [?], though adapting texture-based tracking for standard frame-rate cameras remains the primary method for accessible sports analytics.

3.4. Solution Approach

To transition from a raw monocular video sequence to metric 3D analytics, the proposed solution employs a sequential, decoupled architecture that isolates the structural extraction of the environment from the kinematic tracking of the ball.

The first phase of the pipeline focuses on resolving the spatial geometry of the scene. Due to severe occlusion and the nature of a bowling alley, the playing surface is localized by independently extracting its distinct geometric boundaries (the foul line, the lateral gutters, and the pin deck) using a combination of directional gradient analysis, spatial proximity heuristics, and robust object-centric regression. The intersection of these boundaries yields four coplanar control points in the image plane. By mapping these homogeneous pixel coordinates to the known metric dimensions of a standard bowling lane, a 2D-to-2D planar homography matrix (H) is computed. This mathematical transformation effectively eliminates perspective distortion, providing a rectified, metric bird’s-eye view of the playing surface without requiring intrinsic camera calibration.

Following environmental rectification, the second phase isolates the kinematic trajectory of the bowling ball. The ball is tracked across the temporal sequence to extract its raw 2D center coordinates and apparent pixel radius. To mitigate the fragmentation of data caused by fast motion and object occlusion, temporal interpolation is applied to generate a continuous, smooth mathematical curve in the image plane.

To reconstruct the 3D trajectory, the ball’s image-plane contact point (found by adding its pixel radius to its vertical center) is projected through the homography matrix to obtain its metric ground coordinates (X, Y). By enforcing the constraint that the ball maintains surface contact with the lane, its known physical radius is simply assigned as the Z-coordinate, yielding the final spatial position (X, Y, R).

Finally, to provide comprehensive sports analytics, the system evaluates the rotational dynamics of the ball. By tracking the cyclical, frame-by-frame spatial displacement of high-contrast surface textures across the ball’s visible hemisphere, the angular velocity is mathematically estimated, yielding the overall spin rate over the course of the trajectory.

4. Methodology

4.1. Lane Detection

Lane detection constitutes a fundamental component of the proposed bowling analysis pipeline. Unlike traditional autonomous driving applications, where lane detection is primarily used for navigation, in this work it serves two main purposes: restricting the search region for bowling ball detection and estimating the lane geometry required for perspective rectification. By estimating the lane boundaries, a planar homography can be computed to rectify the image plane into a top-down representation of the bowling lane. This rectification, combined with the assumption that the bowling ball remains in contact with the lane surface and preserves a constant radius, enables the estimation of the ball trajectory in three-dimensional space.

Bowling lane detection presents several challenges, including varying illumination conditions, strong perspective effects, the presence of multiple adjacent lanes with similar geometric structures, and partial occlusions caused by bowling pins. Furthermore, not all lane boundaries exhibit the same visual characteristics. While some borders appear as strong visible edges, others are partially occluded or difficult to identify directly. For this reason, a boundary-specific detection strategy was adopted, where each lane border is estimated using a tailored approach while maintaining a shared preprocessing pipeline. Since the camera remains fixed throughout each bowling recording, the lane geometry is assumed to remain constant across frames. Consequently, lane boundaries are estimated only once using the first frame of the video, and the resulting geometric configuration is reused for the subsequent rectification and trajectory reconstruction stages. This assumption substantially reduces computational cost, since the camera remains fixed throughout the recording.

The proposed lane detection framework consists of three main stages: image preprocessing, edge extraction, and geometric line estimation. Standard preprocessing operations, including color space transformation and Gaussian smoothing, are first applied to enhance relevant structural information and suppress noise. Subsequently, edge detection and Hough-based line extraction are employed, with specialized heuristics introduced depending on the boundary being estimated.

4.1.1 Single-channel image transformation

Since edge detection methods operate on single-channel intensity images, an initial transformation from the RGB image space $R^{n \times m \times 3}$ to a single-channel representation $R^{n \times m}$ is required. The effectiveness of this transformation directly affects the separability of lane boundaries from the surrounding scene, influencing both edge detection robustness and the sensitivity of subsequent line extraction methods. The objective of this stage is to maximize the contrast between relevant lane structures and the background while minimizing the effect of illumination variability across videos. Bowling recordings may differ considerably in lighting conditions, reflections on the lane surface, and camera exposure settings, making the choice of transformation particularly important.

To identify the optimal preprocessing step, several distinct single-channel conversion methods were evaluated. As a baseline, the standard OpenCV grayscale conversion was tested, which computes a weighted combination of the RGB channels according to perceptual luminance [? ?]

$$Y = 0.299 R + 0.587 G + 0.114 B \quad (1)$$

where R, G, and B represent the red, green, and blue image channels, respectively.

An alternative standard structural representation, lightness [?], was also tested to observe if averaging the maximum and minimum color components provided better separability for the lane textures. The lightness

channel from the HLS color space is mathematically defined as:

$$L = \frac{\max(R, G, B) + \min(R, G, B)}{2}$$

Beyond standard conversions, domain-specific approaches were explored. A color transformation inspired by lane detection approaches in the autonomous driving domain [?], calculated as

$$I = R + G - B \quad (2)$$

was evaluated. This transformation emphasizes red and green channel intensities while suppressing the blue channel, theoretically increasing the contrast between bright lane markings and darker road-like surfaces.

To further improve robustness against illumination changes, Principal Component Analysis (PCA) [?] was applied. By projecting the RGB channels into a lower-dimensional space and retaining only the first principal component, this method aims to isolate the dominant structural features of the lane regardless of specific lighting conditions.

Finally, to assess whether boundary transitions could be isolated strictly through chromaticity rather than intensity, the hue channel from the HSV color space was evaluated.

To determine the most robust approach for the final pipeline, the five single-channel transformations were systematically evaluated against various edge detectors (Sobel[?], Canny[?], and Laplacian[?]) to measure their impact on boundary localization. Preliminary evaluations indicated that the default grayscale conversion, when paired with a Sobel filter, provided the most robust and consistent boundary localization across the varying illumination conditions of the dataset. The detailed quantitative analysis and discussion of this evaluation are presented in RESULTS SECTION(CITAR). Consequently, the default grayscale transformation paired with Sobel edge detection was selected as the foundational preprocessing step for the lane detection architecture.

4.1.2 Image smoothing

Before edge detection, Gaussian blurring is applied to reduce high-frequency noise and suppress small intensity variations that may generate false edge responses. In bowling lane imagery, reflections on the lane surface and texture variations can introduce spurious gradients that negatively affect boundary detection. A Gaussian filter with kernel size 5×5 is applied to the single-channel transformed image. This preprocessing step improves the stability of subsequent edge detection methods while preserving the dominant structural information of the lane boundaries.

4.1.3 Edge detection

After image smoothing, edge detection is applied to extract candidate structures corresponding to lane boundaries. Three classical gradient-based and intensity-based methods were evaluated: Sobel, Canny, and Laplacian filters. These methods differ in how they estimate intensity discontinuities and in their sensitivity to noise and texture variations.

The Sobel operator computes first-order image gradients and is used to approximate horizontal and vertical intensity changes. It is particularly useful for emphasizing directional structures, making it suitable for detecting lane boundaries with a strong geometric orientation.

The Canny detector follows a multi-stage process involving gradient computation, non-maximum suppression, and hysteresis thresholding. This approach typically produces thinner and more connected edges, but its performance is sensitive to threshold selection and illumination variability.

The Laplacian operator is a second-order derivative method that detects regions of rapid intensity change regardless of direction. While it can highlight fine structures, it is also more sensitive to noise and may produce fragmented edge maps in low-contrast regions. All three methods were implemented and evaluated across multiple bowling videos. Based on qualitative performance and stability of the resulting edge maps, the Sobel operator was selected as the primary edge detector for the proposed pipeline. Depending on the boundary being estimated, the gradient direction is adapted to emphasize either horizontal or vertical structures, improving the extraction of relevant lane edges prior to line fitting.

4.1.4 Line extraction using the Hough transform

Following edge detection, candidate line segments corresponding to lane boundaries are extracted using the probabilistic Hough transform [?]. The Hough transform [?] is particularly suitable for lane detection tasks, as it enables robust identification of line structures even when edges are fragmented or partially incomplete. The method operates by mapping edge points from the image space into a parametric representation of lines. Since infinitely many lines may pass through a given image point, each edge pixel votes for all compatible line

hypotheses in the parameter space. Line candidates receiving strong consensus from multiple edge points accumulate higher votes and are identified as probable geometric structures. In this work, the probabilistic Hough transform (HoughLinesP) is employed. Unlike the standard Hough transform, which returns infinite lines represented in polar form (ρ, θ) , the probabilistic formulation directly estimates finite line segments defined by their endpoints (x_1, y_1, x_2, y_2) . This representation is particularly advantageous for the proposed pipeline, as subsequent processing relies on geometric properties such as segment orientation, intersections, and proximity to expected lane locations. Additionally, the probabilistic formulation reduces computational cost by operating on a subset of edge points while maintaining robust detection performance. The resulting line candidates are further filtered according to geometric constraints consistent with the expected lane structure. In particular, line orientation is used to reject candidates that deviate substantially from the expected direction of the corresponding lane boundary. Since multiple plausible line hypotheses may still remain, additional boundary-specific heuristics are introduced in later stages of the pipeline to determine the final lane borders.

4.1.5 Bottom boundary detection

The bottom boundary of the bowling lane corresponds to the nearest and most visually distinctive lane edge in the scene. Unlike other lane borders, this boundary is typically less affected by occlusion and exhibits a strong horizontal structure, making it the most reliable reference for initial lane geometry estimation. To improve robustness and reduce interference from irrelevant scene content, a spatial cropping is first applied to focus on the lower-central region of the frame, where the bottom boundary is expected to appear. This restriction reduces noise from distant lane regions and peripheral elements. The cropped region is then converted into a single-channel representation and smoothed using Gaussian filtering. Edge detection is performed using a Sobel operator configured to emphasize horizontal intensity changes, which are most consistent with the orientation of the bottom lane boundary. After edge extraction, the probabilistic Hough transform is applied to identify candidate line segments. Since multiple line hypotheses may be generated due to reflections and lane texture, the resulting segments are filtered based on geometric constraints. In particular, only lines with slopes within a predefined threshold are retained, ensuring consistency with the expected near-horizontal orientation of the bottom boundary. Among the remaining candidates, the most stable line is selected as the bottom boundary estimate.

4.1.6 Lateral boundary detection

The lateral boundaries of the bowling lane present a significantly more challenging detection problem compared to the bottom boundary. Unlike the single dominant horizontal edge observed at the bottom of the lane, the lateral region contains multiple parallel lane markings across adjacent lanes. This introduces ambiguity, as several valid line candidates may be detected in the image, and the correct pair of boundaries depends on the specific lane being analyzed. To address this, the same preprocessing pipeline is applied, consisting of single-channel transformation, Gaussian smoothing, and edge detection using a Sobel operator configured to emphasize vertical structures. Unlike the bottom boundary case, no spatial cropping is applied, as relevant lane structures may extend across a larger portion of the frame.

After edge extraction, the probabilistic Hough transform is used to obtain candidate line segments. A preliminary geometric filter is applied to discard near-horizontal segments, ensuring that only approximately vertical structures are retained. Since multiple valid vertical line hypotheses typically remain, an additional selection strategy is required. To disambiguate between adjacent lanes, a reference point within the target lane is introduced. In this work, the image center is used as a default reference under the assumption that the central lane is the target of interest. However, this point can be modified to track alternative lanes if required. A horizontal line passing through this reference point is then considered, and intersections between this line and all detected line candidates are computed. Each candidate line is evaluated based on the horizontal distance between its intersection point and the reference point. The left and right lane boundaries are selected as the closest valid lines on each side of the reference point, respectively. This heuristic effectively resolves the ambiguity introduced by multiple parallel lane structures by converting the problem into a geometric proximity selection task in one dimension.

4.1.7 Top boundary detection

The top boundary of the bowling lane presents a fundamentally different challenge compared to the other lane borders, as it is frequently occluded by bowling pins. This occlusion breaks the continuity of the lane marking and makes classical edge-based or line-based methods such as the Hough transform unreliable. Additionally, feature-based approaches such as Scale Invariant Feature Transform (SIFT)[?] were considered but found to be ineffective due to the small size of the pins in the image and the limited number of stable keypoints they generate at the typical camera distance.

To address these limitations, a template matching-based approach is adopted, leveraging prior knowledge of the approximate appearance and scale of bowling pins. The method searches for occurrences of a pin template within the upper region of the frame, where pins are expected to appear. Since the apparent size of pins varies across videos due to perspective and camera distance, a multi-scale search strategy is applied. The template is resized across a range of scales, and matching is performed at each scale independently. This increases robustness to variations in pin size and improves detection consistency across different recordings.

After matching, multiple overlapping detections are typically produced for the same physical pin. To remove redundant detections, non-maximum suppression is applied, retaining only the most confident bounding boxes. From each remaining detection, representative keypoints are extracted, corresponding either to the center or the lower portion of the bounding box, depending on the selected configuration.

Given that template matching can still produce false positives in visually similar regions, a robust geometric fitting stage is applied. The extracted keypoints are first filtered using Random Sample Consensus (RANSAC)[?] regression to remove outliers. A final first-order polynomial is then fitted to the inlier points using least squares, resulting in a continuous estimate of the top lane boundary.

4.1.8 Lane boundary integration and corner estimation

Once the bottom, lateral, and top lane boundaries are estimated independently, they are combined to recover a consistent geometric representation of the bowling lane. Since each boundary is represented as a line segment in image space, the lane corners can be obtained by computing the pairwise intersections between adjacent boundaries.

Specifically, the bottom and top boundaries are intersected with the left and right lateral boundaries to obtain the four lane corners: bottom-left, bottom-right, top-right, and top-left. These intersection points define a quadrilateral that approximates the perspective projection of the bowling lane in the image.

4.2. Ball Tracking

The products derived from the bowling release are the ball’s 2D trajectory along the lane and its spin. The trajectory is defined by localizing the geometric center of the ball through the clip’s frames, whereas the parameters of spin are axis orientation and angular velocity. The previously obtained data is used to define the region of interest in which we’ll find the ball’s projection as a circle in every frame. Once the trajectory data is available, the ball’s locations across consecutive frames are used to determine the positional change of its features and to derive its spin parameters. With both products in hand, we can then generate a visualization that consistently follows the ball’s motion and proceed with the 3D reconstruction of the scene.

4.2.1 Trajectory detection

The ball-tracking workflow consists of three main steps: video preprocessing, ball-candidate detection, and trajectory computation. We begin by isolating the scene’s region of interest and applying noise-reduction and contrast-enhancement filters to each frame. Then, the *Circle Hough Transform* (CHT) algorithm is applied to generate a dataset of all geometries that could represent the ball’s position in each frame. Finally, we select the geometries that best form a coherent trajectory from the candidates, interpolate any missing values in the set, and smooth the final trajectory. In the end, we obtain a ball-center dataset with the image geometric coordinates and their corresponding radii.

Firstly, each frame must be manipulated to reduce video noise so that the CHT consistently detects the ball. This process starts by masking all pixels outside the ROI using the coordinates retrieved in the previous section ???. Then, median blurring is applied to remove high-frequency noise from light reflections in the lane, followed by adaptive histogram equalization to increase contrast specifically in the region where the ball is located. This last step is particularly helpful because the ROI is expected to have a relatively consistent surface across the frames. We end up with a black-and-white clip in which the ball is clearly distinguishable from the rest of the scene, and the lane reflections are smoothed out.

Once we have the preprocessed clip, we apply OpenCV’s built-in CHT algorithm. The library implements the *Hough gradient method*, which first applies *Canny edge detection* and then computes the local gradient for every nonzero pixel via the first-order Sobel x and y derivatives. These gradients are used to populate the two-dimensional *accumulator plane* (??), in which potential circle centers are tracked. For every edge pixel, a line is projected along its gradient from a specified minimum to a maximum distance, incrementing the accumulator cells along the way. The local maxima in this plane are selected as candidate centers. Finally, for each center, the algorithm evaluates the distances to nearby edge pixels to determine the optimal radius, preserving the circle only if it receives sufficient support and maintains a minimum distance from previously validated centers.

The result is a dataset containing every detected circle, its x- and y-center coordinates, and its radius for each frame.

This algorithm efficiently detects circles across video frames; however, it may also introduce noise due to numerical instabilities [?], leading to false detections. The third step of the workflow focuses on filtering all circles representing the ball in each frame, resulting in a final defined trajectory. The methodology for discernment was inspired by [?], in which they elaborate a weighted directed graph to determine the ball’s path in a soccer broadcast. The difference in our implementation is that a bowling release is less dynamic than a soccer kick; its behavior is more predictable, and its trajectory is easier to generalize. The main idea is that each ball candidate in a frame is represented by a graph node, which is then directionally connected to all other candidates in the five subsequent frames. This is done because the previous step might miss the ball in consecutive frames. In the end, the longest path in the graph is retrieved, which represents the actual ball trajectory. To fill in the missing frames in the trajectory, cubic spline interpolation is applied to both the x- and y-center coordinates, followed by a light Gaussian filter ($\sigma = 5$) to eliminate noise, and a simple quadratic least squares interpolation for the radius values.

The entire workflow yields the final dataset, which contains the coordinates and radius of each circle from the first to the last detected frame of the clip. This data will be crucial for the following steps in spin detection and 3D reconstruction.

4.3. Lane rectification and 3D Reconstruction

Once both lane geometry and ball trajectory have been estimated in the image plane, the next step consists of transforming the scene into a physically meaningful coordinate system aligned with the real-world dimensions of the bowling lane. Due to perspective distortion, measurements in the image space are not directly interpretable in metric terms. For this reason, a perspective rectification is performed using a homography grounded in the known geometry of the lane.

4.3.1 Homography-based lane rectification

Given the four lane corner points obtained from the intersection of the estimated boundaries, a planar homography is computed to map points from the image plane to a top-down representation of the bowling lane. Unlike a generic projective transformation, this mapping is constrained by the known physical dimensions of a standard bowling lane, namely a width of 1.066 m and a length of 18.29 m [?]. The destination coordinate system is defined directly in metric units, such that the rectified plane corresponds to a physically scaled representation of the lane surface. This establishes a direct correspondence between image coordinates and real-world measurements. The homography H is estimated such that:

$$\mathbf{x}_{\text{metric}} = H \mathbf{x}_{\text{image}}$$

where $\mathbf{x}_{\text{metric}}$ lies in a metric coordinate system expressed in meters. As a result, the rectified output represents the lane in a perspective-corrected and metrically scaled coordinate system.

4.3.2 Ball trajectory transformation and 3D reconstruction

The reconstructed trajectory is obtained through a sequence of geometric transformations that map image-space detections into a 3D representation expressed in the coordinate system of the bowling lane.

First, the ball center is detected in the image plane as a sequence of positions over time, together with the corresponding estimated radii. The ball-lane contact point is then computed under the assumption that the ball remains in continuous contact with the lane surface, by correcting the detected center using the estimated radius along the vertical image axis.

These contact points are projected into the metric lane coordinate system using the previously computed homography, which is defined from the known physical dimensions of the lane. This yields a 2D trajectory expressed in real-world units on the lane plane.

Finally, a 3D representation of the motion is obtained by lifting the contact trajectory along the vertical axis by a distance equal to the ball radius, resulting in a set of 3D coordinates expressed in the lane reference frame. The final output of this stage is the trajectory of the ball expressed in the 3D coordinate system of the lane. The resulting representation is consistent with both the geometric constraints of the lane and the physical dimensions of the ball [?].

5. Results

6. Results

7. Experimental Activity and Analysis of Experimental Results

7.1. Preprocessing and Edge Detection comparative analysis

To empirically validate the optimal preprocessing and feature extraction strategy for the lane detection pipeline, a comparative analysis was conducted across eight distinct bowling video clips. These clips were specifically selected to encompass a variety of challenging real-world conditions, including varying camera exposures, surface glare, and heterogeneous lighting. The study evaluated the combinatorial performance of five single-channel transformations (Default Grayscale, Lightness, Hue, $R + G - B$, and PCA) paired with three standard edge detection algorithms (Sobel, Canny, and Laplacian). A trial was deemed successful strictly if the resulting gradient map enabled the subsequent Hough Transform and heuristic filters to correctly segment the entire playing surface by accurately isolating all four geometric boundaries (the foul line, the pin deck, and both lateral gutters). The quantitative success rates for each combination are detailed in Table 1.

An analysis of the aggregated results reveals clear performance disparities not only among the color transformations but, crucially, among the edge detection algorithms. The Sobel operator emerged as the unequivocally strongest edge detection method for correctly identifying lane borders, outperforming Canny and Laplacian across almost every grayscale transformation. In contrast, while the multi-stage Canny detector and the second-order Laplacian filter are exceptionally strong edge detectors, this extreme sensitivity proved detrimental in a visually cluttered bowling alley. Qualitative analysis of the failed clips revealed that Canny and Laplacian frequently detected crisp, high-contrast edges entirely outside the target playing surface, such as architectural lines on the background wall, adjacent lane gutters, or the bowler’s silhouette. These highly defined but irrelevant structures generated overwhelmingly strong candidate lines in the Hough space. Consequently, these false positives bypassed the spatial and geometric heuristics of the pipeline, leading to incorrect boundary selection and severely deformed segmentation quadrilaterals (see Figure 1). Sobel, by providing a more subdued and directionally targeted gradient map, allowed the geometric filters to consistently lock onto the true lane boundaries.

When observing the single-channel transformations paired with this dominant Sobel filter, a divergence between general cross-filter robustness and peak specific performance became apparent. In terms of general stability across all detectors, the domain-specific $R + G - B$ transformation proved highly robust. It successfully identified boundaries in 5, 4, and 3 clips when paired with the Sobel, Canny, and Laplacian filters, respectively. This consistency validates the underlying hypothesis inherited from autonomous driving literature: mathematically amplifying the red and green channels while suppressing the blue channel effectively isolates the warm, wooden tones of the bowling lane from the cooler, heavily shadowed regions of the gutters and background.

Despite the baseline stability of the $R + G - B$ transformation, the highest success rate was achieved by the standard Default Grayscale conversion paired with the Sobel filter, which successfully localized boundaries in 6 out of the 8 clips. The perceptual luminance weighting of the default grayscale conversion provided the cleanest directional intensity gradients at the lane edges when acted upon by the Sobel operator. Consequently, optimizing for the highest peak accuracy across diverse video inputs, this specific combination was selected as the final implementation for the pipeline.

Finally, the analysis provided valuable insights into the theoretical limitations of chromatic edge detection. It was initially hypothesized that the Hue channel from the HSV color space might yield interesting results by isolating pure color transitions and entirely negating the impact of absolute brightness and shadows. In practice, however, the Hue channel demonstrated the poorest performance of the study, completely failing to localize boundaries across almost all filters. Qualitative analysis of the output (see Figure 2) revealed the cause of this failure: the Hue representation effectively erased the critical geometric boundaries (particularly the bottom foul line) from the image entirely. This occurs because the physical demarcations of a bowling lane (such as a dark foul line or shaded gutters contrasting against the light wooden surface) are primarily defined by sharp shifts in luminance and intensity, rather than distinct shifts in core color. By stripping away the intensity information, the Hue channel eliminates the necessary gradient contrast, causing the borders to blend indistinguishably into the background. This visual evidence conclusively confirms that robust boundary extraction in a bowling setting relies fundamentally on structural intensity rather than pure chromaticity.

Figure 1: Examples of heuristic failure due to over-sensitive edge detection. Left: Canny detects background wall structures. Right: Laplacian detects the bowler’s leg and adjacent lane. Both result in an incorrect homography quadrilateral.

Table 1: Quantitative success rates of boundary detection for each preprocessing and edge detection combination across 8 video clips.

Grayscale Method	Edge Method	clip1	clip2	clip3	clip4	clip5	clip6	clip7	clip8	Sum
default	sobel	1	1	1	0	1	1	1	0	6
default	canny	0	0	1	0	1	0	0	0	2
default	laplacian	0	0	1	0	1	0	0	0	2
lightness	sobel	1	0	1	0	1	1	1	0	5
lightness	canny	1	0	1	0	0	1	0	0	3
lightness	laplacian	0	0	1	0	0	1	0	0	2
hue	sobel	0	1	0	0	0	0	0	0	1
hue	canny	0	0	0	0	0	0	0	0	0
hue	laplacian	0	1	1	0	0	0	0	0	2
r_g_minus_b	sobel	1	1	0	0	1	1	1	0	5
r_g_minus_b	canny	1	1	1	0	1	0	0	0	4
r_g_minus_b	laplacian	1	1	1	0	0	0	0	0	3
pca	sobel	1	0	1	0	1	1	1	0	5
pca	canny	1	0	1	0	0	1	0	0	3
pca	laplacian	0	0	1	0	1	0	1	0	3

Figure 2: Output of the Hue channel on the lower lane segment. By discarding luminance data, the critical intensity contrast of the foul line is completely erased, resulting in a failure to detect the bottom boundary.

7.2. Top Boundary Localization via Object-Centric Regression

Before finalizing the multi-scale template matching approach, feature-based object detection methods, specifically the SIFT, were empirically evaluated for pin localization. While SIFT is generally highly robust to scale and rotation variations, it proved ineffective in this specific environment. The primary limitation arises from the small relative size of the pins within the global image frame at typical camera distances. Furthermore, because bowling pins possess a smooth, uniform surface with minimal internal texture, they generate an insufficient number of stable, high-contrast keypoints. Qualitative analysis of the feature matching outputs (see Figure 3) revealed that the algorithm consistently failed to localize the target pins on the deck. Instead, the limited keypoints extracted from the pin template produced spurious, incorrect matches with high-frequency noise and visually complex elements in the background, such as wall graphics. Due to this inability to reliably detect the target objects, feature-based extraction was discarded in favor of the intensity-based template matching approach, which successfully leverages the macro-structural silhouette of the pins rather than internal textural features.

As established in the methodology, the top boundary of the lane (the pin deck interface) presents a unique geometric challenge: the physical edge is heavily occluded by the bowling pins themselves, making standard gradient-based edge detection ineffective. To overcome this occlusion, an object-centric approach was evaluated, employing multi-scale template matching to detect the pins and utilizing the bottom-center coordinates of their bounding boxes as spatial candidate points for the boundary. The specific selection of the bottom-center coordinate, rather than the geometric centroid (midpoint) of the bounding box, is a critical design choice. Empirical analysis of the bounding box outputs revealed that using the centroid artificially elevated the regression line, mapping it across the midsections of the pins rather than the lane surface. Because the physical objective is to map the pin deck interface, the bottom-center coordinate provides a geometrically accurate anchor directly to the true floor plane.

While multi-scale template matching successfully localized the target pins, the raw output was inherently noisy. As anticipated in a visually complex bowling environment, the algorithm frequently generated false positive detections, predominantly from pin reflections on the highly polished synthetic oil and visually similar architectural features in the background. Interestingly, the template matching also frequently detected pins actively standing in adjacent parallel lanes. Rather than treating these adjacent detections as noise, the pipeline exploits them to increase the stability of the boundary estimation. Because bowling alleys are constructed with coplanar, horizontally aligned pin decks, pins in adjacent lanes share the exact same geometric baseline as

the target lane. Consequently, these adjacent detections serve as valid, extended data points that drastically improve the stability of the linear estimation.

To robustly extract the true boundary from this spatial data while filtering out the actual noise, a RANSAC regressor was evaluated. RANSAC is uniquely suited for this task due to its outlier rejection capabilities. By iteratively estimating a first-degree polynomial model from random subsets of the spatial coordinates, the algorithm successfully identified the dominant linear structure (the continuous base of the pins extending across the visible lanes) while ignoring coordinate data from reflections and background architecture that did not conform to this geometry.

Qualitative analysis of the regression output (see Figure 4) demonstrates the high efficacy of this approach across varying camera angles and lane designs. The RANSAC algorithm successfully segmented the raw candidate points into a strict inlier consensus set (visualized as blue points) and a rejected outlier set (visualized as red points). The inlier set actively incorporates valid pins from both the target and adjacent lanes, exploiting the extended horizontal baseline to anchor the regression more firmly. The red spatial outliers correctly encapsulate the false detections caused by reflections and background noise. By fitting the final line strictly to these fortified blue inliers, the pipeline generated a highly accurate and stable top boundary line (visualized in green). This confirms that decoupling the structural detection into an object-centric model effectively bypasses severe visual occlusions at the pin deck.

Figure 3: Failure of SIFT feature matching for pin detection. Due to the small scale and lack of internal texture on the pins, the algorithm fails to find valid corresponding keypoints on the pin deck. Instead, it generates spurious matches between the template (top left) and random high-frequency background elements.

Figure 4: Performance of the RANSAC regressor on multi-scale template matching output. The bottom-center coordinates of all detected bounding boxes are plotted. The algorithm successfully classifies true pins (including those actively exploited on adjacent, aligned lanes) as inliers (blue dots). It effectively rejects false positives from reflections and background structures as outliers (red dots). The final fitted polynomial (green line) accurately maps the occluded pin deck boundary using the extended geometric baseline.

7.3. Ball Tracking (Agrega lo que quieras)

7.4. Perspective Rectification and Homography Validation

Following the successful segmentation of the lane boundaries and the robust localization of the pin deck via RANSAC regression, the four intersecting corners of the playing surface (foul line and pin deck intersecting with the lateral gutters) are extracted to form a 2D quadrilateral in the image plane. To translate these distorted pixel coordinates into a uniform spatial environment, a planar homography matrix is computed. This projective transformation maps the source quadrilateral onto a normalized, top-down rectangular destination plane representing the physical proportions of the bowling lane.

To qualitatively validate the accuracy of the computed homography matrices, the original perspective video frames were warped to generate a direct bird’s-eye view of the lane. As demonstrated in Figure 5, this transformation successfully rectifies the physical surface across various video clips characterized by differing camera positions, focal lengths, and lighting conditions. The geometric efficacy of the homography is visually confirmed by the structural alignment within these warped images: the longitudinal wood grain and target arrows maintain strict verticality without tapering toward a vanishing point. Because even minor pixel localization errors in the preceding boundary detection modules would severely distort this projection, the undistorted spatial structure achieved validates the stability of the foundational geometric model. Consequently, these homography matrices provide a reliable mathematical foundation for the subsequent mapping of the balls’ real-world trajectories.

Figure 5: Bird’s-eye view perspective rectification across multiple test videos. The computed homography matrices successfully eliminate perspective distortion, rendering the wood boards perfectly parallel. This confirms the high precision of the foundational boundary detection algorithms.

7.5. Trajectory Reconstruction and 3D Spatial Visualization

Building upon the validated homography matrix and the continuously interpolated 2D tracking data, the final objective of the pipeline is to reconstruct the bowling ball's trajectory in real-world dimensions. As shown in Figure 6, the interpolated trajectory is first visually rendered as a continuous trace directly onto the perspective video feed. Simultaneously, these continuous pixel coordinates are multiplied by the homography matrix to project the ball's positional path onto a two-dimensional bird's-eye view. It is important to note that, for graphical visualization purposes, the rendered top-down lane shown on the right side of the figure is intentionally distorted, compressed longitudinally and expanded laterally compared to true physical proportions. While the underlying homographic projection calculates metrically accurate spatial coordinates, mapping the projected trajectory onto this visually scaled lane significantly improves readability, exaggerating the transverse dynamics to make the ball's lateral hook clearly perceptible to the viewer.

To provide a comprehensive spatial analysis, this metrically accurate 2D planar data is subsequently extended into a full three-dimensional physical representation. The pipeline defines a real-world coordinate system anchored at the bottom-left corner of the physical lane (the origin), with the X and Y axes corresponding to the standardized physical length and width of the playing surface. The Z axis elevates the physical radius of the bowling ball above the floor plane, allowing for a volumetric representation of the movement.

Qualitative analysis of the final spatial visualizations (see Figure 7) confirms the success of the end-to-end pipeline. Qualitative analysis of the 3D plots confirms the success of the end-to-end pipeline. The reconstructed trajectory exhibits a smooth curvature that accurately mirrors the ball's hook dynamics from the original video. Ultimately, this proper spatial alignment demonstrates that the integrated computer vision pipeline successfully translates raw 2D pixel data into a reliable 3D kinematic model.

Figure 6: 2D Trajectory mapping. Left: The interpolated tracking data is traced in red on the original perspective image plane. Right: The pixel coordinates are warped via the computed homography matrix into a normalized bird's-eye view. The graphical representation of the lane is intentionally scaled to exaggerate transverse movement, revealing the true lateral curvature of the throw.

Figure 7: Three-dimensional spatial reconstruction of the lane and ball trajectory. The 2D homographic data is mapped into a physical coordinate space, rendering the lane's true dimensions, the origin point, and a volumetric model of the ball following the calculated trajectory path.

8. Equations

This section gives some examples of writing mathematical equations in your thesis. Maxwell's equations read:

$$\left\{ \begin{array}{l} \nabla \cdot \mathbf{D} = \rho, \\ \nabla \times \mathbf{E} + \frac{\partial \mathbf{B}}{\partial t} = \mathbf{0}, \\ \nabla \cdot \mathbf{B} = 0, \\ \nabla \times \mathbf{H} - \frac{\partial \mathbf{D}}{\partial t} = \mathbf{J}. \end{array} \right. \quad \begin{array}{l} (3a) \\ (3b) \\ (3c) \\ (3d) \end{array}$$

Equation (3) is automatically labeled by `cleveref`, as well as Equation (3a) and Equation (3c). Thanks to the `cleveref` package, there is no need to use `\eqref`. Equations have to be numbered only if they are referenced in the text.

Equations (4), (5), (6), and (7) show again Maxwell's equations without brace:

$$\nabla \cdot \mathbf{D} = \rho, \quad (4)$$

$$\nabla \times \mathbf{E} + \frac{\partial \mathbf{B}}{\partial t} = \mathbf{0}, \quad (5)$$

$$\nabla \cdot \mathbf{B} = 0, \quad (6)$$

$$\nabla \times \mathbf{H} - \frac{\partial \mathbf{D}}{\partial t} = \mathbf{J}. \quad (7)$$

Equation (8) is the same as before, but with just one label:

$$\left\{ \begin{array}{l} \nabla \cdot \mathbf{D} = \rho, \\ \nabla \times \mathbf{E} + \frac{\partial \mathbf{B}}{\partial t} = \mathbf{0}, \\ \nabla \cdot \mathbf{B} = 0, \\ \nabla \times \mathbf{H} - \frac{\partial \mathbf{D}}{\partial t} = \mathbf{J}. \end{array} \right. \quad (8)$$

9. Figures, Tables and Algorithms

Figures, Tables and Algorithms have to contain a Caption that describes their content, and have to be properly referred in the text.

9.1. Figures

For including pictures in your text you can use `TikZ` for high-quality hand-made figures [?], or just include them with the command

```
\includegraphics[options]{filename.xxx}
```

Here xxx is the correct format, e.g. `.png`, `.jpg`, `.eps`,



Figure 8: Caption of the Figure.

9.2. Tables

Within the environments `table` and `tabular` you can create very fancy tables as the one shown in Table 2.

Example of Table (optional)

	column1	column2	column3
row1	1	2	3
row2	α	β	γ
row3	alpha	beta	gamma

Table 2: Caption of the Table.

You can also consider to highlight selected columns or rows in order to make tables more readable. Moreover, with the use of `table*` and the option `bp` it is possible to align them at the bottom of the page. One example is presented in Table 3.

9.3. Algorithms

Pseudo-algorithms can be written in $\text{\LaTeX}$ with the `algorithm` and `algorithmic` packages. An example is shown in Algorithm 1.

Algorithm 1 Name of the Algorithm

```
1: Initial instructions
2: for for – condition do
3:   Some instructions
4:   if if – condition then
5:     Some other instructions
6:   end if
7: end for
8: while while – condition do
9:   Some further instructions
10: end while
11: Final instructions
```

10. Some further useful suggestions

Theorems have to be formatted as follows:

Theorem 10.1. *Write here your theorem.*

Proof. If useful you can report here the proof.

Propositions have to be formatted as follows:

Proposition 10.1. *Write here your proposition.*

How to insert itemized lists:

- first item;
- second item.

How to write numbered lists:

1. first item;
2. second item.

11. Use of copyrighted material

Each student is responsible for obtaining copyright permissions, if necessary, to include published material in the thesis. This applies typically to third-party material published by someone else.

12. Plagiarism

You have to be sure to respect the rules on Copyright and avoid an involuntary plagiarism. It is allowed to take other persons' ideas only if the author and his original work are clearly mentioned. As stated in the Code of Ethics and Conduct, Politecnico di Milano *promotes the integrity of research, condemns manipulation and the infringement of intellectual property*, and gives opportunity to all those who carry out research activities to have an adequate training on ethical conduct and integrity while doing research. To be sure to respect the copyright rules, read the guides on Copyright legislation and citation styles available at:

<https://www.biblio.polimi.it/en/tools/courses-and-tutorials>

	column1	column2	column3	column4	column5	column6
row1	1	2	3	4	5	6
row2	a	b	c	d	e	f
row3	α	β	γ	δ	ϕ	ω
row4	alpha	beta	gamma	delta	phi	omega

Table 3: Highlighting the columns

You can also attend the courses which are periodically organized on "Bibliographic citations and bibliography management".

13. Conclusions

A final section containing the main conclusions of your research/study and possible future developments of your work have to be inserted in the section "Conclusions".

14. Bibliography and citations

Your thesis must contain a suitable Bibliography which lists all the sources consulted on developing the work. The list of references is placed at the end of the manuscript after the chapter containing the conclusions. It is suggested to use the BibTeX package and save the bibliographic references in the file `bibliography.bib`. This is indeed a database containing all the information about the references. To cite in your manuscript, use the `\cite{}` command as follows:

Here is how you cite bibliography entries: [?], or multiple ones at once: [? ?].

The bibliography and list of references are generated automatically by running BibTeX [?].

A. Appendix A

If you need to include an appendix to support the research in your thesis, you can place it at the end of the manuscript. An appendix contains supplementary material (figures, tables, data, codes, mathematical proofs, surveys, ...) which supplement the main results contained in the previous sections.

B. Appendix B

It may be necessary to include another appendix to better organize the presentation of supplementary material.

Abstract in lingua italiana

Qui va l'Abstract in lingua italiana della tesi seguito dalla lista di parole chiave.

Parole chiave: qui, le parole chiave, della tesi, in italiano

Acknowledgements

Here you might want to acknowledge someone.